

PLATFORM WHITEPAPER 2025

Sovereign AI Infrastructure with Carbon-Aware Scheduling

A comprehensive technical and economic analysis of the Harch Platform — five GPU hubs, 1,798 accelerators, 81.5% renewable energy, and 89% lower carbon intensity than industry average.

PUBLISHED: MARCH 2025

VERSION: 1.0

CLASSIFICATION: PUBLIC

Table of Contents

Table of Contents	1
Executive Summary	3
The AI Compute Challenge	4
GPU Scarcity and the Access Crisis	4
The Carbon Crisis in AI Compute	4
The Sovereignty Gap	4
Introducing the Harch Platform	6
Architecture Overview	6
HarchOS: The Carbon-Aware Orchestrator	6
Five GPU Hubs Deep-Dive	7
Ouarzazate Hub	7
Dakhla Hub	7
Benguerir Hub	7
Tanger Hub	8
Casablanca Hub	8
Carbon-Aware Scheduling	9
How Carbon-Aware Scheduling Works	9
Real-Time Routing and Deferral	9
Carbon Intensity: Harch vs. Hyperscalers	9
Submarine Connectivity	10
Cable Systems and Latency	10
Sovereign Security Framework	11
Physical Security	11
Cybersecurity Architecture	11
Data Residency and Compliance	11
AI Platform Services	13
Training and Fine-Tuning	13

Inference at Scale 13

MLOps and SDK 13

Sustainable Compute Economics 14

 Total Cost of Ownership Advantage 14

 Pricing Model 14

Competitive Landscape 16

Get Started 17

 Request a Custom Quote 17

 Schedule a Technical Briefing 17

 Join the Team 17

 Connect With Us 17

Executive Summary

The global artificial intelligence industry is experiencing an unprecedented surge in demand for high-performance compute infrastructure. From large language models with hundreds of billions of parameters to real-time inference serving millions of users, the computational requirements of modern AI workloads have outpaced the capacity of traditional cloud providers. Organizations face a triple crisis: GPU scarcity that delays training runs by weeks, cost structures that make sustained AI development prohibitively expensive, and a carbon footprint that conflicts with increasingly stringent regulatory mandates and corporate sustainability commitments.

Harch Corp presents the Harch Platform, a purpose-built AI infrastructure platform that fundamentally reimagines how GPU compute is provisioned, scheduled, and consumed. Anchored by five strategically located GPU hubs across the Kingdom of Morocco, the platform delivers 1,798 GPUs today with a 500 MW deployment pipeline scheduled through 2027-2028. What distinguishes the Harch Platform is its carbon-aware scheduling engine, which routes workloads in real time to the hub with the lowest carbon intensity at any given hour, achieving an average grid carbon intensity of approximately 47 gCO₂/kWh, which is 89% below the industry average of roughly 450 gCO₂/kWh.

With 81.5% average renewable energy penetration across its hubs, submarine cable connectivity delivering 8 ms latency to European financial centers, a sovereign security framework compliant with EU and African data residency regulations, and a 99.99% uptime SLA backed by hardware-level redundancy, the Harch Platform offers a compelling alternative to hyperscaler GPU clouds. The \$1.14 billion investment program underpinning this infrastructure delivers a projected 24.7% internal rate of return, validating both the commercial viability and the strategic importance of sovereign, sustainable AI compute for Europe, Africa, and the Middle East.

The AI Compute Challenge

GPU Scarcity and the Access Crisis

The explosive growth of foundation models has created a structural deficit in GPU availability. NVIDIA reported that demand for H100-class accelerators exceeded supply by a factor of three throughout 2023 and 2024, with wait times extending to 36 weeks for large cluster reservations. This scarcity disproportionately impacts organizations outside the hyperscaler ecosystem, including national research institutes, mid-market AI companies, and sovereign entities that cannot compete with the procurement power of the largest cloud providers. The result is a two-tier market in which a handful of incumbents control access to the most critical resource in the AI economy.

Even when GPU capacity is available, the cost of sustained training and inference on hyperscaler platforms remains prohibitive for many workloads. Training a single 70-billion-parameter model on demand can cost upwards of \$2 million, while inference at scale compounds these expenses month after month. For organizations in Europe and Africa, these economics are further distorted by currency fluctuations, cross-border data transfer fees, and the absence of regional pricing parity programs.

The Carbon Crisis in AI Compute

The environmental impact of AI compute has emerged as a defining challenge for the industry. A single large-scale training run can emit as much carbon as five automobiles over their entire lifetimes, and the aggregate carbon footprint of global AI infrastructure now exceeds that of several small nations. Hyperscaler data centers, predominantly located in regions with fossil-fuel-heavy grids, operate at average carbon intensities of 400-550 gCO₂/kWh. As the EU Corporate Sustainability Reporting Directive (CSRD) and similar regulations take effect, organizations face mandatory carbon disclosure requirements that make the location and energy source of AI compute a strategic concern, not merely an ethical one.

The Sovereignty Gap

European and African organizations face a unique set of constraints that hyperscaler architectures were not designed to address. Data sovereignty regulations, including the EU AI Act, GDPR data residency requirements, and emerging African data protection frameworks, demand that certain categories of data be processed within defined jurisdictions. The current concentration of GPU capacity in US-based availability zones creates a structural sovereignty gap that cannot be closed by contractual mechanisms alone. Physical infrastructure, operated under local jurisdiction, with auditable security controls, is the only path to genuine sovereign AI compute.

Challenge Dimension	Current State	Impact
GPU Availability	3x demand-supply gap	36-week wait times for H100 clusters

Training Cost (70B model)	\$2M+ per run	Prohibitive for non-hyperscalers
Hyperscaler Carbon Intensity	400-550 gCO ₂ /kWh	CSRD non-compliance risk
EU Data Residency	No GPU hubs in compliant zones	Sovereignty gap for EU/Africa
Energy Cost (EU avg)	\$0.064/kWh	2-4x cost of renewable-rich regions

Table 1: Key dimensions of the AI compute challenge facing organizations today.

Introducing the Harch Platform

Architecture Overview

The Harch Platform is a vertically integrated AI infrastructure system that spans the full stack from physical data center operations to cloud-native platform services. Unlike general-purpose cloud providers that retrofit GPU capacity into existing server fleets, every element of the Harch Platform, from power procurement and cooling design to network fabric and orchestration software, is optimized for AI workloads. This purpose-built approach eliminates the performance penalties and operational friction that arise when GPU computing is layered onto infrastructure designed for web-scale applications.

The platform architecture is organized into four layers: the Physical Layer, comprising five GPU hubs with purpose-built facilities optimized for high-density compute; the Network Layer, anchored by four submarine cable systems providing redundant, low-latency connectivity to Europe, Africa, and the Middle East; the Orchestration Layer, powered by HarchOS, a custom Kubernetes-based scheduler with carbon-aware routing; and the Services Layer, offering training, fine-tuning, inference, and MLOps capabilities through a unified API and SDK. This layered architecture enables each tier to be independently scaled and maintained while maintaining strict end-to-end quality of service guarantees.

HarchOS: The Carbon-Aware Orchestrator

At the heart of the Harch Platform sits HarchOS, a custom-built orchestration system that extends Kubernetes with real-time carbon-intensity-aware scheduling. HarchOS continuously ingests grid carbon intensity data from each hub location, combines it with workload requirements and service-level objectives, and makes placement decisions that minimize carbon emissions without compromising performance or availability. When a training job is submitted, HarchOS evaluates the current carbon intensity at each hub, the available GPU capacity, the network latency to the data source, and any data residency constraints before selecting the optimal execution site.

For time-insensitive workloads such as batch training, HarchOS can defer execution to periods of peak renewable generation, further reducing the carbon footprint by 15-30% compared to immediate scheduling. For latency-sensitive inference workloads, HarchOS ensures placement at the hub with the best network path to the end user while still preferring lower-carbon options when latency requirements permit. This intelligent scheduling approach is what transforms renewable energy potential from a static attribute into a dynamic operational advantage.

Layer	Components	Key Capability
Physical	5 GPU hubs, 1,798 GPUs, 500 MW pipeline	Purpose-built high-density AI compute
Network	4 submarine cables, redundant backbone	8 ms latency to Europe, 99.99% SLA
Orchestration	HarchOS, carbon-aware scheduler, K8s	Real-time carbon-intensity routing

Services	Training, fine-tuning, inference, MLOps	Unified API/SDK, auto-scaling
----------	---	-------------------------------

Table 2: Harch Platform four-layer architecture summary.

SECTION 04

Five GPU Hubs Deep-Dive

The Harch Platform operates across five GPU hubs strategically distributed throughout the Kingdom of Morocco, each selected for its unique combination of renewable energy resources, network proximity to target markets, and climatic advantages for efficient cooling. Together, these hubs provide 1,798 GPUs today, with the 500 MW deployment pipeline expanding capacity to serve the full spectrum of AI workloads from training to inference at scale.

Ouarzazate Hub

The flagship Ouarzazate hub is located adjacent to the Noor Ouarzazate Solar Complex, the largest concentrated solar power facility in the world, with over 580 MW of installed solar capacity. This proximity provides the hub with a direct renewable energy pipeline that delivers 97.2% renewable electricity, resulting in a grid carbon intensity of just 18 gCO₂/kWh, among the lowest of any GPU compute facility globally. The hub currently operates 800 GPUs and serves as the primary training cluster for large-scale foundation models. The arid Saharan climate enables free-air cooling for over 8,000 hours per year, dramatically reducing the energy overhead associated with traditional chiller-based cooling systems. The Ouarzazate hub is designed for expansion under the 500 MW pipeline, with plans to scale to multi-thousand GPU capacity by 2028.

Dakhla Hub

The Dakhla hub is positioned in the southern coastal region where the combination of Atlantic wind corridors and high solar irradiance creates one of the most favorable renewable energy environments on the planet. Operating 400 GPUs with 94.8% renewable energy penetration and a carbon intensity of 32 gCO₂/kWh, Dakhla is particularly notable for its energy cost advantage: at \$0.018/kWh, electricity costs are 72% below the European average of \$0.064/kWh. This cost advantage translates directly into lower per-GPU-hour pricing for Harch customers. The hub also benefits from consistent Atlantic breezes that provide natural cooling, further reducing operational energy consumption.

Benguerir Hub

Located in the Mohammed VI Polytechnic University green campus, the Benguerir hub operates 350 GPUs with 88.5% renewable energy and a carbon intensity of 55 gCO₂/kWh. This hub serves as the primary research and development cluster, benefitting from direct proximity to Morocco's leading AI research institutions and a pipeline of technical talent. The Benguerir hub is designed for fine-tuning and experimentation workloads, with direct high-bandwidth connectivity to the Ouarzazate training cluster for seamless model transfer. The university campus environment also supports pilot programs for emerging cooling technologies, including direct-to-chip liquid cooling and immersion cooling systems that are expected to reduce PUE below 1.05 in future

deployments.

Tanger Hub

The Tanger hub, operating 200 GPUs with 82.1% renewable energy and a carbon intensity of 95 gCO₂/kWh, is the network gateway to Europe. Located just 14 km from the Spanish coast at the Strait of Gibraltar, Tanger provides the lowest-latency path for inference workloads serving European end users. The hub connects directly to all four submarine cable systems, making it the primary ingress and egress point for traffic between the Harch Platform and European networks. Tanger's role as a Mediterranean logistics hub also provides excellent fiber connectivity to West Africa and the Middle East, supporting Harch's multi-continental reach.

Casablanca Hub

The Casablanca hub operates 48 GPUs with 45% renewable energy and a carbon intensity of 210 gCO₂/kWh. While its renewable fraction is lower than the other hubs, Casablanca serves a critical function as the enterprise services hub, providing low-latency access to Morocco's financial district and business centers. This hub handles API gateway services, control plane operations, and compliance-sensitive workloads that must remain within Moroccan jurisdiction under the CNDP data protection framework. Casablanca also hosts the platform's Network Operations Center (NOC) and Security Operations Center (SOC), providing 24/7 monitoring and incident response across all five hubs.

Hub	GPUs	Renewable %	Carbon (gCO ₂ /kWh)	Energy Cost (\$/kWh)	Primary Role
Ouarzazate	800	97.2%	18	\$0.022	Large-scale training
Dakhla	400	94.8%	32	\$0.018	Cost-optimized compute
Benguerir	350	88.5%	55	\$0.025	R&D; and fine-tuning
Tanger	200	82.1%	95	\$0.030	European inference gateway
Casablanca	48	45%	210	\$0.038	Enterprise and compliance

Table 3: Comparative overview of the five Harch GPU hubs.

Carbon-Aware Scheduling

How Carbon-Aware Scheduling Works

Carbon-aware scheduling is the cornerstone of the Harch Platform's sustainability advantage. Unlike conventional cloud schedulers that optimize exclusively for cost, latency, or utilization, HarchOS incorporates real-time carbon intensity data as a first-class scheduling dimension. The system continuously monitors the carbon intensity of the electrical grid at each hub location through integrations with Morocco's ONEE (Office National de l'Electricite et de l'Eau Potable) grid operator data feeds and European Entso-E transparency platform APIs. These data streams are supplemented by on-site solar irradiance and wind speed sensors that predict renewable generation output for the next 24-48 hours with high accuracy.

When a workload is submitted to the platform, HarchOS evaluates a multi-objective cost function that weights carbon intensity, GPU availability, data residency constraints, network latency, and cost. The relative weights are configurable per workload: a training job for a sustainability-focused research institution might weight carbon intensity at 60%, while a real-time inference service for a financial trading platform might weight latency at 70%. This flexibility ensures that carbon optimization never comes at the expense of performance requirements that are critical to the customer's business.

Real-Time Routing and Deferral

For batch training workloads that have flexible scheduling windows, HarchOS implements a deferral mechanism that shifts execution to periods of peak renewable generation. Analysis of historical grid data shows that solar generation in the Ouarzazate region peaks between 10:00 and 15:00 local time, during which carbon intensity drops to as low as 8 gCO₂/kWh. By aligning compute-intensive workloads with these windows, the platform achieves a 15-30% reduction in operational carbon emissions compared to immediate scheduling, with no impact on training completion time when the deferral window is shorter than the total slack in the project schedule.

For inference workloads that cannot be deferred, HarchOS performs real-time geographic routing. When an inference request arrives from a European client, the system evaluates the current carbon intensity at both the Tanger hub (closest geographically) and the Ouarzazate hub (lowest carbon intensity). If the latency budget permits the additional 3-4 ms of network delay to Ouarzazate, and the carbon intensity differential is significant, the request is routed to the lower-carbon hub. This dynamic routing ensures that sustainability is optimized continuously, not merely as a static attribute of the data center's location.

Carbon Intensity: Harch vs. Hyperscalers

The carbon intensity advantage of the Harch Platform is substantial and systematic. The average grid carbon intensity across all five hubs is approximately 47 gCO₂/kWh, which represents an 89% reduction compared to the industry average of approximately 450 gCO₂/kWh. This is not a marginal improvement but a categorical shift: Harch workloads operate on a fundamentally different carbon curve than those running on conventional

cloud infrastructure. For organizations subject to mandatory carbon reporting, this difference can transform AI compute from a carbon liability into a carbon advantage.

Provider / Region	Avg Carbon (gCO ₂ /kWh)	Renewable %	vs. Harch Avg
Harch Platform (avg)	47	81.5%	Baseline
Harch Ouarzazate	18	97.2%	-62%
AWS eu-west-1 (Ireland)	316	42%	+572%
GCP europe-west1 (Belgium)	178	56%	+279%
Azure West Europe (Netherlands)	334	38%	+611%
Industry Average	450	~30%	+857%

Table 4: Carbon intensity comparison between Harch hubs and major hyperscaler regions.

SECTION 06

Submarine Connectivity

The Harch Platform's strategic value as a bridge between European, African, and Middle Eastern markets is underpinned by its connectivity infrastructure. Four submarine cable systems provide redundant, high-bandwidth paths between Morocco and key global internet exchanges, ensuring that the platform can deliver the low-latency, high-throughput connectivity that AI workloads demand.

Cable Systems and Latency

The primary connectivity asset is the Morocco-Spain cable system, which provides a direct path from Tanger to Spain with a round-trip latency of approximately 8 ms to major European peering points in Madrid, Marseille, and Frankfurt. This latency is comparable to intra-European paths and enables Harch to serve European inference workloads with performance indistinguishable from EU-based data centers. The Africa Coast to Europe (ACE) cable provides a second Mediterranean path with connectivity to France and Portugal, while the SAT-3/WASC cable links Morocco to West African markets including Senegal, Nigeria, and South Africa. The Maroc Telecom Cable provides additional Mediterranean redundancy.

Redundancy is engineered at every layer: each cable system lands at multiple points of presence, traffic is load-balanced across cables using BGP anycast, and failover between cable systems is automatic with sub-second convergence. This architecture ensures that the failure of any single cable system, however unlikely, does not result in service degradation. The 99.99% SLA commitment is directly enabled by this multi-path connectivity design.

Cable System	Route	RTT Latency to EU	Capacity	Redundancy Role
--------------	-------	-------------------	----------	-----------------

Morocco-Spain	Tanger to Spain	8 ms	Multiple Tbps	Primary EU path
ACE	Morocco to France/Portugal	12 ms	Multiple Tbps	Secondary EU path
SAT-3/WASC	Morocco to West Africa	35-80 ms	Multiple Tbps	Africa connectivity
Maroc Telecom	Mediterranean ring	10 ms	Multiple Tbps	EU tertiary path

Table 5: Submarine cable systems connecting the Harch Platform to global markets.

SECTION 07

Sovereign Security Framework

Physical Security

Each Harch GPU hub is designed to meet or exceed the physical security standards required for Tier IV data center certification. Facilities are built within secured perimeters featuring 24/7 armed security, biometric access controls, CCTV surveillance with 90-day retention, and anti-tailgating mantraps at all entry points. The Ouarzazate and Dakhla hubs benefit from their remote locations, which provide natural physical security advantages, while the Tanger and Casablanca hubs employ additional urban-hardening measures including vehicle barriers, blast-resistant facades, and electromagnetic shielding. All hubs maintain physical separation between customer compute zones and shared infrastructure, ensuring that no unauthorized physical access to customer hardware is possible under any circumstances.

Cybersecurity Architecture

The Harch Platform implements a zero-trust security model in which every network flow, whether between hubs, between customer tenants, or between platform services, is authenticated and encrypted. All data in transit is protected by TLS 1.3 with mutual authentication, while data at rest is encrypted using AES-256 with customer-managed keys stored in Hardware Security Modules (HSMs). The platform's network architecture enforces strict micro-segmentation: each customer tenant operates within an isolated network fabric with dedicated virtual firewalls, and all inter-tenant traffic is blocked by default. East-west traffic within the platform is monitored by an inline intrusion detection system that inspects every packet for indicators of compromise.

The Security Operations Center (SOC) in Casablanca provides 24/7 monitoring across all five hubs, with automated threat detection, incident triage, and response orchestration powered by a combination of commercial SIEM platforms and custom machine-learning models trained on the platform's unique traffic patterns. The SOC maintains partnerships with national CERT teams in Morocco, the EU, and participating African nations to ensure coordinated response to cross-border threats. Regular penetration testing by independent third-party firms validates the effectiveness of security controls, with results shared with customers under NDA.

Data Residency and Compliance

Data residency is a non-negotiable requirement for many organizations, particularly those in regulated industries such as financial services, healthcare, and government. The Harch Platform provides configurable data

residency controls at the hub, region, and country level, ensuring that customer data never leaves the designated jurisdiction without explicit authorization. For EU customers, workloads can be configured to execute exclusively within the Tanger hub, which maintains compliance with GDPR data processing requirements. For Moroccan and African customers, the Casablanca hub provides compliance with the CNDP (Commission Nationale de controle de la Protection des Donnees a caractere Personnel) framework.

The platform's compliance program is aligned with ISO 27001, SOC 2 Type II, and the EU Cloud Code of Conduct. Annual audits by Big Four accounting firms validate the effectiveness of controls, and audit reports are available to customers under appropriate confidentiality agreements. Harch is also pursuing FedRAMP authorization for US government workloads and is engaging with the African Union's emerging data governance framework to ensure compliance as continental regulations evolve.

AI Platform Services

Training and Fine-Tuning

The Harch Platform provides a comprehensive suite of AI training services designed to support the full lifecycle of model development. The distributed training engine supports data parallelism, tensor parallelism, and pipeline parallelism across multi-node GPU clusters, enabling efficient scaling from single-node experimentation to thousand-GPU training runs. The platform natively supports all major training frameworks including PyTorch, JAX, and TensorFlow, with pre-configured environments that include optimized CUDA libraries, NCCL collective operations, and distributed checkpointing.

Fine-tuning services build on the training infrastructure with additional optimizations for parameter-efficient methods. LoRA, QLoRA, and full fine-tuning workflows are supported through a declarative configuration interface that abstracts away infrastructure complexity. Customers specify the base model, the dataset, the fine-tuning method, and the desired output characteristics; the platform handles GPU allocation, data preprocessing, training orchestration, evaluation, and model packaging. This approach reduces the time from dataset to deployed model from weeks to hours for common fine-tuning workflows.

Inference at Scale

Production inference on the Harch Platform is powered by a high-throughput serving engine that supports auto-scaling from zero to thousands of GPU instances in response to demand. The inference engine implements continuous batching, speculative decoding, and PagedAttention to maximize GPU utilization and minimize per-request latency. For latency-critical applications, the platform supports dedicated GPU reservations that guarantee capacity and response times, even during peak demand periods.

The inference service is accessible through OpenAI-compatible API endpoints, enabling customers to migrate existing applications with minimal code changes. Model deployment is streamlined through a one-click deployment pipeline that handles model optimization (including quantization and compilation with TensorRT and vLLM), canary rollouts, A/B testing, and automated rollback. Comprehensive observability, including request-level latency percentiles, throughput metrics, and GPU utilization, is available through a real-time dashboard and Prometheus-compatible metrics endpoint.

MLOps and SDK

The Harch Platform SDK provides native client libraries for Python, TypeScript, and Go, enabling developers to integrate Harch services directly into their existing ML pipelines and applications. The SDK abstracts away infrastructure concerns such as GPU allocation, data transfer, and job scheduling, allowing developers to focus on model development and deployment. The MLOps toolkit includes experiment tracking, model registry, feature store integration, and CI/CD pipelines for automated model retraining and deployment. These tools are designed to integrate seamlessly with popular open-source MLOps platforms including MLflow, Weights and Biases, and Kubeflow, ensuring that customers are not locked into proprietary tooling.

Sustainable Compute Economics

Total Cost of Ownership Advantage

The economic case for the Harch Platform rests on three pillars: significantly lower energy costs, reduced cooling overhead, and the elimination of cross-border data transfer fees for European and African customers. At the Dakhla hub, electricity costs of \$0.018/kWh represent a 72% reduction compared to the European average of \$0.064/kWh. For a 1,000-GPU training cluster operating at 300 kW per node, this energy cost advantage alone translates to annual savings exceeding \$12 million compared to running the same workload in a European data center. When combined with free-air cooling savings of 20-30% on PUE-related costs, the total operational expenditure advantage reaches 40-55% for sustained compute workloads.

Capital expenditure is also favorable. Morocco's competitive construction and labor costs, combined with government incentives for renewable energy investments, result in data center build costs that are 30-40% lower than equivalent facilities in Western Europe. These savings are reflected in the platform's pricing model, which offers GPU-hour rates that are competitive with, and in many cases lower than, the major hyperscaler providers, despite operating at a smaller scale. The economic model is validated by the \$1.14 billion investment program's projected 24.7% internal rate of return, which demonstrates that sovereign, sustainable AI compute is not only environmentally responsible but commercially superior.

Pricing Model

The Harch Platform offers a flexible pricing model designed to accommodate the diverse needs of AI organizations, from startups running their first experiments to enterprises deploying production inference at scale. On-demand pricing provides immediate access to GPU capacity with per-second billing and no commitment. Reserved pricing offers significant discounts of 40-60% for 1-3 year commitments, with guaranteed capacity and priority scheduling. Spot pricing provides access to unused capacity at discounts of up to 70% off on-demand rates, with 30-second termination notices that enable checkpointing for fault-tolerant workloads.

Carbon credits are integrated directly into the pricing model: customers receive verified carbon offsets proportional to the emissions avoided by running workloads on the Harch Platform rather than on conventional cloud infrastructure. These offsets are certified under the Verified Carbon Standard (VCS) and can be applied toward the customer's Scope 2 emissions reporting, creating a tangible financial benefit beyond the direct cost savings on compute.

Cost Dimension	Harch Platform	EU Hyperscaler	Savings
Energy (\$/kWh)	\$0.018-0.038	\$0.055-0.080	45-72%
PUE (effective)	1.05-1.15	1.20-1.40	12-25%
Data Transfer (EU)	Included	\$0.05-0.12/GB	100%

Carbon Cost (est.)	\$0 (offsets generated)	\$15-40/MWh equiv.	100%
GPU-hour (reserved, H100)	\$1.80-2.20	\$2.80-3.50	30-40%

Table 6: Total cost of ownership comparison between Harch Platform and EU-based hyperscaler alternatives.

Competitive Landscape

The AI infrastructure market is dominated by three hyperscaler platforms (AWS, Google Cloud, and Microsoft Azure) that collectively control over 65% of global GPU cloud capacity. OVHcloud and Scaleway represent European alternatives, but their GPU offerings remain limited in scale and lack the purpose-built optimization that AI workloads demand. The Harch Platform occupies a unique position in this landscape: it is the only provider that combines sovereign jurisdiction with carbon-aware scheduling, renewable energy economics, and low-latency connectivity to European and African markets.

The table below provides a detailed comparison across the dimensions that matter most to organizations making strategic infrastructure decisions: carbon intensity, renewable energy, latency to European users, sovereign data residency, and pricing competitiveness. On every dimension, the Harch Platform delivers a differentiated advantage that cannot be replicated by retrofitting existing cloud infrastructure.

Dimension	Harch	AWS	GCP	Azure	OVHcloud
Carbon (gCO ₂ /kWh)	47	~400	~320	~380	~250
Renewable Energy	81.5%	~42%	~56%	~48%	~35%
EU Latency	8 ms	5-15 ms	3-12 ms	4-14 ms	2-8 ms
Sovereign Residency	Yes (Morocco)	Limited	Limited	Limited	EU only
Carbon-Aware Sched.	Native	No	Partial	No	No
GPU-hour (H100 res.)	\$1.80-2.20	\$2.80-3.50	\$2.90-3.60	\$3.00-3.70	\$2.50-3.20
Africa Connectivity	Direct	Indirect	Indirect	Indirect	Minimal
SLA	99.99%	99.99%	99.99%	99.99%	99.99%

Table 7: Competitive comparison across key infrastructure dimensions.

Get Started

The Harch Platform is available to organizations of all sizes, from individual researchers and startups to large enterprises and government agencies. Getting started is straightforward and designed to minimize time-to-value, with multiple engagement pathways tailored to different organizational needs and readiness levels.

Request a Custom Quote

For organizations with defined AI workload requirements, the Harch sales team provides custom pricing proposals that account for workload characteristics, commitment duration, data residency needs, and sustainability targets. Quotes are delivered within 48 hours and include a detailed carbon impact assessment showing the emissions avoided by choosing the Harch Platform over conventional cloud alternatives. Contact the sales team at sales@harch.co to initiate the process.

Schedule a Technical Briefing

Technical briefings provide a 60-minute deep dive into the platform architecture, carbon-aware scheduling capabilities, and service offerings, conducted by a senior solutions architect. Briefings can be tailored to specific use cases, including large-scale training, production inference, fine-tuning workflows, or sovereign compliance requirements. Schedule a briefing at briefing@harch.co or through the platform portal at portal.harch.co.

Join the Team

Harch Corp is hiring across engineering, operations, and commercial roles. As a rapidly scaling infrastructure platform at the intersection of AI and sustainability, Harch offers unique career opportunities for professionals who want to shape the future of sustainable AI compute. Open positions include GPU systems engineers, platform reliability engineers, carbon-intensity data scientists, solutions architects, and business development managers. Visit careers.harch.co for current openings and application details.

Connect With Us

- Email: info@harch.co
- Platform Portal: portal.harch.co
- Documentation: docs.harch.co
- LinkedIn: linkedin.com/company/harch-corp
- GitHub: github.com/harch-corp